

This is the final polish. The resulting architecture is now robust against drift, immutability bugs, and schema confusion.

Here are the finalized documents. Hand these to the ck3raven-dev agent.

Document 1: The Architecture Specification

Save as: docs/arch/CANONICAL_REPLY_SYSTEM.md

Canonical Architecture: Centralized Reply System

Status: APPROVED STANDARD

Target: Global (MCP, Core Library, Policy)

Goal: Eliminate "Context Poisoning" and "Silent Failures" via deterministic, registry-based replies.

1. Core Philosophy: "Replies, Not Errors"

We replace ad-hoc Python exceptions with a structured **Reply System**.

- **Exceptions** are for control flow *within* a function.
- **Replies** are for communication across boundaries (Tool \$\to\$ Agent).

The Razor-Sharp Separation of Concerns

We strictly define who owns which outcome to prevent logic leakage.

Component	Responsibility	Allowed Replies	Forbidden Replies
WorldAdapter	Physics. "Does this path exist in the lens?"	S, I	D (Denied)
Enforcement	Governance. "Is this permitted by contract?"	S, D	I (Invalid)

Infrastructure	Transport. "Did the system crash?"	E	D (Denied)
-----------------------	---	---	------------

2. The Reply Taxonomy

The reply_type discriminator determines the Agent's **Next Action**.

Type	Name	Owner	Semantics	Agent Reaction
S	Success	All	Operation completed.	Continue.
I	Invalid	WorldAdapter / Tool Boundary	Request cannot be executed as stated (bad path, bad input, schema violation).	Self-Correct: Fix input format/path.
D	Denied	Enforcement	Valid request rejected by Policy/Governance.	Escalate: Request Token or Change Scope.
E	Error	Infrastructure	System failure, unhandled exception, timeout.	Stop: Report trace_id. Do not retry.

3. The Registry Contract (Anti-Drift)

To prevent message drift, we use a **Single Source of Truth**. Replies are defined by a code and a message_key.

Compatibility Contract: Agents must branch on reply_type and code. **NEVER**

branch on message text.

3.1 The Schema

Python

```
@dataclass(frozen=True)
class Reply:
    reply_type: Literal["S", "I", "D", "E"]
    code: str          # Canonical ID (e.g., WA-RES-I-001)
    message_key: str   # Registry key (e.g., PATH_NOT_FOUND)
    params: Dict[str, Any]  # Data for template rendering
    data: Optional[Dict] = None # Machine-readable payload (decisions, diffs)
    trace_id: Optional[str] = None

    @property
    def message(self) -> str:
        # Rendered from Registry (Convenience Only)
        return REGISTRY[self.code].format(**self.params)
```

3.2 Canonical Code Format

Format: LAYER-AREA-TYPE-NNN

- **Layer:** WA (World), EN (Enforcement), MCP (Transport), CT (Contract).
- **Area:** RES (Resolution), IO, READ, WRITE, LINT, GATE, SYS.
- **Type:** S, I, D, E.

Registry Examples:

- WA-RES-I-001: "Path '{path}' does not exist."
- EN-WRITE-D-002: "Write denied to '{path}'. Outside Contract Scope."
- MCP-SYS-S-900: "Legacy tool returned raw payload; wrapped as Success."

4. Implementation Architecture: The Safety Wrapper

The wrapper is a **pure safety net**. It guarantees the transport contract and immutability.

Responsibilities:

1. **Trace ID:** Generate unique ID for every request.

2. **Immutability:** Never mutate the result; return a new instance with the trace ID.
3. **Auto-Wrap:** If tool returns legacy data, wrap as Reply(S) using code MCP-SYS-S-900.
4. **Catch-All:** Catch Exception \$\to\$ Log Stack \$\to\$ Return Reply(E).

Pseudo-Code:

Python

```
def mcp_safe_tool(func):
    def wrapper(*args, **kwargs):
        trace_id = generate_uuid()
        try:
            result = func(*args, **kwargs)

            # Legacy Support (Phase 1 Only)
            if not isinstance(result, Reply):
                return Reply(
                    reply_type="S",
                    code="MCP-SYS-S-900",
                    message_key="LEGACY_WRAP",
                    params={},
                    data=result,
                    trace_id=trace_id
                )

            # Immutability: Return new instance with trace_id
            return dataclasses.replace(result, trace_id=trace_id)

        except Exception as e:
            # 1. Capture Stack Trace to Disk
            log_exception_to_disk(trace_id, e)
            # 2. Return Clean 'E' Reply
            return Reply(
                reply_type="E",
                code="MCP-SYS-E-001",
                message_key="SYS_CRASH",
                params={"err": str(e)},
                trace_id=trace_id
            )
    return wrapper
```

Document 2: The Implementation Plan

Save as: docs/plans/REPLY_SYSTEM_IMPLEMENTATION.md

Canonical Reply System Implementation Plan

Status: APPROVED

Goal: Secure the transport immediately, then refactor core logic.

Phase 1: Infrastructure & Safety Net (Immediate)

Objective: Stop crashes and "silent" failures. Introduce the Registry.

1. **Create Core Infrastructure:**
 - src/ck3raven/core/replies.py: Define Reply dataclass (frozen).
 - src/ck3raven/core/reply_registry.py: Define REGISTRY dict.
 - *Must include:* MCP-SYS-S-900 ("Legacy tool returned raw payload").
2. **Implement Safety Wrapper:**
 - tools/ck3lens_mcp/safety.py: Implement @mcp_safe_tool.
 - *Logic:* Use dataclasses.replace for trace ID injection (Immutability).
3. **Wrap Tool Entrypoints:**
 - Decorate every function in tools/ck3lens_mcp/server.py.
 - *Verification:* Verify existing tools still work (wrapped with MCP-SYS-S-900).

Phase 2: Domain Adoption (The Refactor)

Objective: Replace ad-hoc exceptions with semantic I and D replies.

1. **Refactor Enforcement (The Guard):**
 - Target: src/ck3raven/policy/Enforcement.py.
 - Action: Change check_write_permission to return Reply(D) instead of raising PermissionError.
 - *Integration:* Update server.py to check if result.reply_type == 'D': return result.
2. **Refactor WorldAdapter (The Lens):**
 - Target: src/ck3raven/core/WorldAdapter.py.
 - Action: Return Reply(I) for missing paths.

- *Constraint:* WorldAdapter **NEVER** returns D.

Phase 3: Strict Mode (Cleanup)

Objective: Enforce the standard.

1. Disable Auto-Wrap:

- Update @mcp_safe_tool to return Reply(E) if the underlying tool returns a raw dict.
- *Prerequisite:* All tools must be migrated to return Reply objects.

2. Audit:

- Ensure no code branches on reply.message. All logic must use reply.code or reply.reply_type.